

C

robert lupton¹

23 september 2010

¹LooChao@gmail.com

Outline

Introduction

C was designed as a system programming language, to remove the necessity of writing operating systems in assembler. It's one of the large family of languages deriving from Algol.

Introduction

C was designed as a system programming language, to remove the necessity of writing operating systems in assembler. It's one of the large family of languages deriving from Algol.

You've seen "hello world" before: `CANNOT INCLUDE FILE src/hello.c`

Introduction

C was designed as a system programming language, to remove the necessity of writing operating systems in assembler. It's one of the large family of languages deriving from Algol.

You've seen "hello world" before: `CANNOT INCLUDE FILE src/hello.c`
This example comes from "The C Programming Language" by Brian Kernighan and Dennis Ritchie ("K&R")

Declarations

All variables must be declared before they can be used:

```
char c = 'a';  
char *s = "I am a string";  
double d = 1;  
float f = 1.0f;  
int i = 101010101;  
long l = 10;  
short s = 10;
```

Declarations

All variables must be declared before they can be used:

```
char c = 'a';  
char *s = "I am a string";  
double d = 1;  
float f = 1.0f;  
int i = 101010101;  
long l = 10;  
short s = 10;
```

“string” is spelt `char *` for reasons that I'll explain in a bit.

Declarations

All variables must be declared before they can be used:

```
char c = 'a';  
char *s = "I am a string";  
double d = 1;  
float f = 1.0f;  
int i = 101010101;  
long l = 10;  
short s = 10;
```

“string” is spelt `char *` for reasons that I’ll explain in a bit. `short` is actually a shorthand for `short int` (`long` means `long int`). You can also define integral types (`char` and `int`) as being `signed` or `unsigned`.

Declarations

All variables must be declared before they can be used:

```
char c = 'a';  
char *s = "I am a string";  
double d = 1;  
float f = 1.0f;  
int i = 101010101;  
long l = 10;  
short s = 10;
```

“string” is spelt `char *` for reasons that I’ll explain in a bit. `short` is actually a shorthand for `short int` (`long` means `long int`). You can also define integral types (`char` and `int`) as being `signed` or `unsigned`. One especially useful qualifier is `const`:

```
const unsigned long bad = 0xdeadbeef00000000;
```

Declarations

All variables must be declared before they can be used:

```
char c = 'a';  
char *s = "I am a string";  
double d = 1;  
float f = 1.0f;  
int i = 101010101;  
long l = 10;  
short s = 10;
```

“string” is spelt `char *` for reasons that I’ll explain in a bit. `short` is actually a shorthand for `short int` (`long` means `long int`). You can also define integral types (`char` and `int`) as being `signed` or `unsigned`. One especially useful qualifier is `const`:

```
const unsigned long bad = 0xdeadbeef00000000;
```

An unqualified `int` is supposed to be the most efficient integral type on your machine, and is what you’d generally use unless there was some reason not to.

Declaring your own types

C allows you to use your own name for a type:

```
typedef unsigned short U16;
```

Declaring your own types

C allows you to use your own name for a type:

```
typedef unsigned short U16;
```

Why would I want to do this?

Declaring your own types

C allows you to use your own name for a type:

```
typedef unsigned short U16;
```

Why would I want to do this?

CCD image data is typically created using a 16-bit A/D converter, so the natural type for a single pixel is a 2-byte integer.

But C doesn't tell me how large an `int` is; I can find out (using `sizeof(int)`) but I don't want to have to change all my declarations when I move to a new system.

Declaring your own types

C allows you to use your own name for a type:

```
typedef unsigned short U16;
```

Why would I want to do this?

CCD image data is typically created using a 16-bit A/D converter, so the natural type for a single pixel is a 2-byte integer.

But C doesn't tell me how large an `int` is; I can find out (using `sizeof(int)`) but I don't want to have to change all my declarations when I move to a new system.

Actually, these days, I could say:

```
#include <stdint.h>
typedef uint16_t U16;
```

but even that only works if the processor actually has an unsigned 16-bit type.

Declarations

You can mix code and declarations:

```
int pwv = 0;                                // Precipitable Water Va
/*
 * Estimate the PWV based on the altitude
 */
...;
/*
 * Given that estimate, find a better value
 */
const int pwv0 = pwv;                       /* initial estimate of p
```

Declarations

You can mix code and declarations:

```
int pwv = 0;                                // Precipitable Water Va
/*
 * Estimate the PWV based on the altitude
 */
...;
/*
 * Given that estimate, find a better value
 */
const int pwv0 = pwv;                       /* initial estimate of p
```

The ability to declare variables when they are first needed means that you can usually initialize them too; when possible, make them **const**.

Scope

A variable's *scope* is the part of the programme it may be referenced from; in C, a variable's scope is the nearest set of braces (`{ }`), a *block*. If it isn't in a block, a variable is visible globally (i.e. it'll show up when you run `nm` on your object file). If this isn't what you wanted, you can:

Scope

A variable's *scope* is the part of the programme it may be referenced from; in C, a variable's scope is the nearest set of braces (`{}`), a *block*. If it isn't in a block, a variable is visible globally (i.e. it'll show up when you run `nm` on your object file). If this isn't what you wanted, you can:

- Move it into a function — remember, global variables should usually be avoided

Scope

A variable's *scope* is the part of the programme it may be referenced from; in C, a variable's scope is the nearest set of braces (`{}`), a *block*. If it isn't in a block, a variable is visible globally (i.e. it'll show up when you run `nm` on your object file). If this isn't what you wanted, you can:

- Move it into a function — remember, global variables should usually be avoided
- Label it `static` which makes it only visible within the file

Scope

A variable's *scope* is the part of the programme it may be referenced from; in C, a variable's scope is the nearest set of braces (`{}`), a *block*. If it isn't in a block, a variable is visible globally (i.e. it'll show up when you run `nm` on your object file). If this isn't what you wanted, you can:

- Move it into a function — remember, global variables should usually be avoided
- Label it `static` which makes it only visible within the file
- Decide that it really *must* be globally visible, and declare it in a header file:

```
extern int nread;      // Number of times I've read from a fi
```

Scope

A variable's *scope* is the part of the programme it may be referenced from; in C, a variable's scope is the nearest set of braces (`{}`), a *block*. If it isn't in a block, a variable is visible globally (i.e. it'll show up when you run `nm` on your object file). If this isn't what you wanted, you can:

- Move it into a function — remember, global variables should usually be avoided
- Label it `static` which makes it only visible within the file
- Decide that it really *must* be globally visible, and declare it in a header file:

```
extern int nread;    // Number of times I've read from a fi
```

It is generally a good idea to declare a variable in as restricted a scope as possible (the “No Globals” rule is a special case of this one).

Control Flow

C has the usual constructs:

- `if (...) { ... } else if (...) { ...} else { ... }`

Control Flow

C has the usual constructs:

- `if (...) { ... } else if (...) { ...} else { ... }`
- `for (...) { ... }`
- `while { ... }`
- `do { ... } while (...);`

Control Flow

C has the usual constructs:

- `if (...) { ... } else if (...) { ...} else { ... }`
- `for (...) { ... }`
- `while { ... }`
- `do { ... } while (...);`
- `switch (...) { case XXX: ... break; ... }`

Control Flow

C has the usual constructs:

- `if (...) { ... } else if (...) { ...} else { ... }`
- `for (...) { ... }`
- `while { ... }`
- `do { ... } while (...);`
- `switch (...) { case XXX: ... break; ... }`
- `break, continue`

Control Flow

C has the usual constructs:

- `if (...) { ... } else if (...) { ...} else { ... }`
- `for (...) { ... }`
- `while { ... }`
- `do { ... } while (...);`
- `switch (...) { case XXX: ... break; ... }`
- `break, continue`
- `goto`

If statements

```
if (i < 9) {  
    printf("Hello");  
} else {  
    printf("Goodbye");  
}  
printf(" world\n");
```

If statements

```
if (i < 9) {  
    printf("Hello");  
} else {  
    printf("Goodbye");  
}  
printf(" world\n");
```

Legally, you can write this as

```
if (i < 9)  
    printf("Hello");  
else  
    printf("Goodbye");
```

```
printf(" world\n");
```

But we don't recommend it;

If statements

```
if (i < 9) {  
    printf("Hello");  
} else {  
    printf("Goodbye");  
}  
printf(" world\n");
```

Legally, you can write this as

```
if (i < 9)  
    printf("Hello");  
else  
    printf("Goodbye");
```

```
printf(" world\n");
```

But we don't recommend it; it's too easy to write

```
if (i < 9)  
    printf("Hello");  
else
```

For loops

CANNOT INCLUDE FILE `src/for.c` This only works with a compiler that supports the C99 standard (`cc --std=c99 ...`)

For loops

CANNOT INCLUDE FILE `src/for.c` This only works with a compiler that supports the C99 standard (`cc --std=c99 ...`) This is almost equivalent to:

```
#include <stdio.h>

int
main()
{
    int i = 0;
    while (i != 10) {
        printf("Hello world\n");
        ++i;
    }

    return 0;
}
```

Switch I

CANNOT INCLUDE FILE `src/printf.c`

Switch

```
$ make printf && ./printf
$ cc -g -Wall -O3 --std=c99    printf.c    -o printf
a
b
c
:
<space>
<integer>
%
<space>
d
o
n
e
<newline>
```

(N.b. The call was `checkFormat("abc: %d%% done\n");`)

Putting that together

CANNOT INCLUDE FILE `src/switch.c`

Subroutines

Here is a function to add two numbers:

Subroutines

Here is a function to add two numbers:
and here is one to multiply them:

Subroutines

Here is a function to add two numbers:
and here is one to multiply them:
So far, just like Fortran

Subroutines

Here is a function to add two numbers:
and here is one to multiply them:
So far, just like Fortran

Subroutines

Here is a function to add two numbers:

and here is one to multiply them:

So far, just like Fortran

That's a little clumsier — in Fortran (or python) you could have said `x**2`. And note the extra include file, **math.h**. In the old days you needed to link `-lm` too, but modern systems seem to be more forgiving.

Recursion

Subroutines may be called recursively — that is, they may call themselves, either directly or indirectly. There is no limit except the capacity of your computer. CANNOT INCLUDE FILE `src/factorial.c`

Recursion

That may not seem very interesting; it's easy enough to write a loop to calculate factorials.

However, consider a routine

```
integrate (float a, float b, float (*func)(float x))2
```

Recursion

That may not seem very interesting; it's easy enough to write a loop to calculate factorials.

However, consider a routine

```
integrate(float a, float b, float (*func)(float x))2
```

If I need to do a double integral, I can say something like:

```
static float yy;                                // current value of y
```

```
static float func(float x) {  
    return sin(x)*cos(yy);  
}
```

```
static float dfunc(float y) {  
    yy = y;                                     // pass y to func  
    return integrate(0, y, func);  
}
```

```
const double ans = integrate(1, 3, dfunc);
```

to calculate

Complete Example

```
#include <stdio.h>
#include <math.h>

double integrate(const float a, const float b, float (*func)(float),
    const int nstep = 1000;           // number of steps
    const float step = (b - a)/nstep;

    double ans = 0.0;
    float x = a;
    for (int i = 0; i != nstep; ++i) {
        ans += func(x);
        x += step;
    }

    return step*ans;
}

static float yy;                                // current value of y
```

Subroutine Arguments

Question: What does this do?

```
void triple(double x) {  
    x *= 3;  
}  
...  
double x = 1;  
triple(x);
```

Subroutine Arguments

Question: What does this do?

```
void triple(double x) {  
    x *= 3;  
}
```

...

```
double x = 1;  
triple(x);
```

Answer: wastes CPU cycles.

The function `triple` is passed a *copy* of `x`, so **nothing** that `triple` does can affect the program that calls it.

Subroutine Arguments

Question: What does this do?

```
void triple(double x) {  
    x *= 3;  
}
```

...

```
double x = 1;  
triple(x);
```

Answer: wastes CPU cycles.

The function `triple` is passed a *copy* of `x`, so **nothing** that `triple` does can affect the program that calls it.

The solution is to pass a *pointer* to `x`:

```
void triple(double @\color{red}*@x) {  
    @\color{red}*@x *= 3;  
}
```

...

```
double x = 1;  
triple(@\color{red}\&x);
```

Subroutine Arguments

Question: What does this do?

```
void triple(double x) {  
    x *= 3;  
}  
...
```

```
double x = 1;  
triple(x);
```

Answer: wastes CPU cycles.

The function `triple` is passed a *copy* of `x`, so **nothing** that `triple` does can affect the program that calls it.

The solution is to pass a *pointer* to `x`:

```
void triple(double @\color{red}*@x) {  
    @\color{red}*@x *= 3;  
}  
...
```

```
double x = 1;  
triple(@\color{red}\&x);
```

Fortran always passes arguments this way: a famous party trick when I


Prototypes

It is critically important that subroutines' callers and callees agree about the number and types of the arguments (the *signature*). C uses a *prototype* to allow the compiler to check; given

```
void triple(double *x) {  
    *x *= 3;  
}
```

we put the prototype `void triple (double *x);` in a header (".h") file and `#include` it in the file that defines `triple` and also whenever we call `triple`.³

This isn't a lot of extra work and soon becomes second nature.

³Most compilers these days will complain if you don't do this. 

Arrays

Arrays are declared and subscripted using `[]`.

```
int ids[10];
```

Arrays

Arrays are declared and subscripted using `[]`.

```
int ids[10];
```

In C99, the dimension need *not* be known at compile time:

```
void foo(const int n) {  
    int ids[n];  
    ...  
    double x = ids[n/2];  
}
```

The index starts at *0* (not *1*, as in Fortran) as it specifies the distance from the start of the array.

```
printf("ID0: %d\n", ids[0]);
```

Arrays

Arrays are declared and subscripted using `[]`.

```
int ids[10];
```

In C99, the dimension need *not* be known at compile time:

```
void foo(const int n) {  
    int ids[n];  
    ...  
    double x = ids[n/2];  
}
```

The index starts at *0* (not *1*, as in Fortran) as it specifies the distance from the start of the array.

```
printf("ID0: %d\n", ids[0]);
```

This may make more sense when we discuss *pointers*.

n-D Arrays

You can also define n-D arrays:

```
U16 data[4096][2048];
```

```
...
```

```
U16 const peak = data[y][x];           // the (x, y)th pixel
```

The data is stored row-by-row so the *last* index increases fastest if we access pixels in the order in which they're stored (unlike Fortran arrays, in which the *first* index varies fastest).

n-D Arrays

You can also define n-D arrays:

```
U16 data[4096][2048];  
...  
U16 const peak = data[y][x];           // the (x, y)th pixel
```

The data is stored row-by-row so the *last* index increases fastest if we access pixels in the order in which they're stored (unlike Fortran arrays, in which the *first* index varies fastest).

n-D arrays aren't as useful as you might think as they are passed to subroutines as *pointers*; for now all you need to know is that you **must** specify all but the first dimension when passing an n-D array to a subroutine:

```
void debias(U16 data[][2048], const int nrow, const int ncol) {  
    ...  
}
```

n-D Arrays

You can also define n-D arrays:

```
U16 data[4096][2048];  
...  
U16 const peak = data[y][x];           // the (x, y)th pixel
```

The data is stored row-by-row so the *last* index increases fastest if we access pixels in the order in which they're stored (unlike Fortran arrays, in which the *first* index varies fastest).

n-D arrays aren't as useful as you might think as they are passed to subroutines as *pointers*; for now all you need to know is that you **must** specify all but the first dimension when passing an n-D array to a subroutine:

```
void debias(U16 data[][2048], const int nrow, const int ncol) {  
    ...  
}
```

(hmm, I seem to have typed 2048 twice; was that a good idea...?)

Operator Precedence

operator	associativity
() [] -> .	left to right
! ~ ++ -- - (type) * & sizeof()	right to left
* / %	left to right
+ -	left to right
<< >>	left to right
< <= > >=	left to right
== !=	left to right
&	left to right
^	left to right
	left to right
&&	left to right
	left to right
?:	right to left
= += -= etc.	right to left
,	left to right

Macros and the C Pre-Processor (*CPP*)

C provides a simple macro processor which can be used to keep *magic numbers* out of your code[3] instead of

```
char configFile[21];                // name of configuration
fgets(stdin, fileName, 20);
```


Macros and the C Pre-Processor (*CPP*)

C provides a simple macro processor which can be used to keep *magic numbers* out of your code[3] instead of

```
char configFile[21];           // name of configuration
fgets(stdin, fileName, 20);
```

you can write

```
#define FILE_LEN 20           // maximum length for fi

char configFile[FILE_LEN + 1]; // name of configuration
fgets(stdin, fileName, FILE_LEN);
```

Macros and the C Pre-Processor (*CPP*)

C provides a simple macro processor which can be used to keep *magic numbers* out of your code[3] instead of

```
char configFile[21];                // name of configuration
fgets(stdin, fileName, 20);
```

you can write

```
#define FILE_LEN 20                // maximum length for fi

char configFile[FILE_LEN + 1];      // name of configuration
fgets(stdin, fileName, FILE_LEN);
```

writing all macros in CAPITALS is a common convention.

Conditional Compilation

In the bad old days, we wrote lots of code like:

```
#if defined(vms)
    return -1;
#elif defined(HAVE_SELECT) && !defined(USE_POLL)
    if(select(ncheck,&mask,(fd_set *)NULL,(fd_set *)NULL,
            (void *)NULL) != 0) {
        ...
    }
#else
    /* Use poll() */
#endif
```

Conditional Compilation

In the bad old days, we wrote lots of code like:

```
#if defined(vms)
    return -1;
#elif defined(HAVE_SELECT) && !defined(USE_POLL)
    if(select(ncheck,&mask,(fd_set *)NULL,(fd_set *)NULL,
            (void *)NULL) != 0) {
        ...
    }
#else
    /* Use poll() */
#endif
```

Fortunately, in these days of Posix (IEEE 1003; ISO/IEC 9945) there is much less need for of this sort of thing.

Conditional Compilation

In the bad old days, we wrote lots of code like:

```
#if defined(vms)
    return -1;
#elif defined(HAVE_SELECT) && !defined(USE_POLL)
    if(select(ncheck,&mask,(fd_set *)NULL,(fd_set *)NULL,
            (void *)NULL) != 0) {
        ...
    }
#else
    /* Use poll() */
#endif
```

Fortunately, in these days of Posix (IEEE 1003; ISO/IEC 9945) there is much less need for of this sort of thing. The (or at least my) main use of the CPP is:

```
#define DEBUG 1
...
#if DEBUG
int niter = 0;
```

Good uses for macros I

One standard use for macros is to prevent header files being parsed more than once

```
#if !defined(GREET_H)
#define GREET_H
/* Lots of stuff that should be processed only once */
#endif
```

This use of a macro is known as an *include guard*.

Good uses for macros II

CANNOT INCLUDE FILE `src/macros.c`

Good uses for macros II

```
CANNOT INCLUDE FILE src/macros.c
```

```
$ make macros && macros 1 2 3
```

```
cc -g -Wall --std=c99 macros.c -o macros
```

```
macros.c:009: Hello world!
```

```
macros.c:010: My name is macros and I was called with 3 argument
```


Good uses for macros II

CANNOT INCLUDE FILE `src/macros.c`

```
$ make macros && macros 1 2 3
```

```
cc -g -Wall --std=c99 macros.c -o macros
```

```
macros.c:009: Hello world!
```

```
macros.c:010: My name is macros and I was called with 3 argument
```

Unfortunately, this code does not conform to the C99 standard (ISO 9899:1999):

```
$ cc -g -Wall --std=c99 --pedantic-errors macros.c -o macros
```

```
macros.c:3:29: error: ISO C does not permit named variadic macro
```

```
macros.c:7:30: error: ISO C99 requires rest arguments to be used
```

Good uses for macros II

A legal version is: `CANNOT INCLUDE FILE src/macros2.c` Note that we were forced to pass at least one argument (hence the `" "` in the first call).

Bad uses for macros I

You can use the CPP to pretend that you're writing Algol, not C

```
#define IF      if(  
#define THEN   ){  
#define ELSE   } else {  
#define ELIF   } else if (  
#define FI     ;}
```

```
IF x == 0 THEN  
    printf("zero\n");  
ELIF x == 1 THEN  
    printf("one\n");  
ELSE  
    printf("many\n");  
FI
```

Bad uses for macros I

You can use the CPP to pretend that you're writing Algol, not C

```
#define IF      if(
#define THEN    ){
#define ELSE    } else {
#define ELIF    } else if (
#define FI      ;}
```

```
IF x == 0 THEN
    printf("zero\n");
ELIF x == 1 THEN
    printf("one\n");
ELSE
    printf("many\n");
FI
```

These macros come from the original Bourne shell source code in Unix Version 7.

Bad uses for macros II

How about

```
#define MAX(a,b) ((a)>(b)?(a):(b))
```

This one also comes from Steve Bourne.

Bad uses for macros II

How about

```
#define MAX(a,b) ((a)>(b)?(a):(b))
```

This one also comes from Steve Bourne.

Why is this bad? Consider

```
const double fg = MAX(funcs(x), gunks(x));
```

Bad uses for macros II

How about

```
#define MAX(a,b)      ((a)>(b)?(a):(b))
```

This one also comes from Steve Bourne.

Why is this bad? Consider

```
const double fg = MAX(funcs(x), gunks(x));
```

```
double funcs(const double x) {  
    static int x = 0.5;  
    x = sin(x);  
    return x;  
}
```

```
static int i = 0;  
double gunks(const double x) {  
    return (x < ++i) ? 5.6 : 5.9;  
}
```

Bad uses for macros II

How about

```
#define MAX(a,b)      ((a)>(b)?(a):(b))
```

This one also comes from Steve Bourne.

Why is this bad? Consider

```
const double fg = MAX(funcs(x), gunks(x));
```

```
double funcs(const double x) {  
    static int x = 0.5;  
    x = sin(x);  
    return x;  
}
```

```
static int i = 0;  
double gunks(const double x) {  
    return (x < ++i) ? 5.6 : 5.9;  
}
```

Both `funcs` and `gunks` are called twice; this is always inefficient, but in this case catastrophic as they maintain internal state.

Bad uses for macros II

How about

```
#define MAX(a,b)      ((a)>(b)?(a):(b))
```

This one also comes from Steve Bourne.

Why is this bad? Consider

```
const double fg = MAX(funcs(x), gunks(x));
```

```
double funcs(const double x) {  
    static int x = 0.5;  
    x = sin(x);  
    return x;  
}
```

```
static int i = 0;  
double gunks(const double x) {  
    return (x < ++i) ? 5.6 : 5.9;  
}
```

Both `funcs` and `gunks` are called twice; this is always inefficient, but in this case catastrophic as they maintain internal state.

Bad uses for macros II

How about

```
#define MAX(a,b)      ((a)>(b)?(a):(b))
```

This one also comes from Steve Bourne.

Why is this bad? Consider

```
const double fg = MAX(funcs(x), gunks(x));
```

```
double funcs(const double x) {  
    static int x = 0.5;  
    x = sin(x);  
    return x;  
}
```

```
static int i = 0;  
double gunks(const double x) {  
    return (x < ++i) ? 5.6 : 5.9;  
}
```

Both `funcs` and `gunks` are called twice; this is always inefficient, but in this case catastrophic as they maintain internal state.

Structs

If you routinely write code like:

```
void printObjects(const int n, const int id[],
                 const float xcen[], const float ycen[],
                 const float flux[]);

#define NOBJECT 1000                                // Maximum number of obj

int id[NOBJECT];                                     // Object IDs
float xcen[NOBJECT];                                 // x-coordinate of centr
float ycen[NOBJECT];                                 // y-coordinate of centr
float flux[NOBJECT];                                 // object's flux

for (int i = 0; i < n; ++i) {
    id[i] = i;
    xcen[i] = ...;
    ycen[i] = ...;
    flux[i] = ...;
}
```

Structs

A cleaner way to write this is:

```
struct Object {  
    int id;                // Object IDs  
    float xcen;            // x-coordinate of center  
    float ycen;            // y-coordinate of center  
};  
  
typedef struct Object Object;  
  
void printObjects(const int n, const Object objs[]);  
  
#define NOBJECT 1000      // Maximum number of objects  
  
Object objs[NOBJECT];     // our objects  
  
for (int i = 0; i < n; ++i) {  
    objs[i].id = i;  
    objs[i].xcen = ...;  
    objs[i].ycen = ...;  
}
```

Structs

A cleaner way to write this is:

```
struct Object {
    int id;           // Object IDs
    float xcen;       // x-coordinate of center
    float ycen;       // y-coordinate of center
};

typedef struct Object Object;

void printObjects(const int n, const Object objs[]);

#define NOBJECT 1000 // Maximum number of objects

Object objs[NOBJECT]; // our objects

for (int i = 0; i < n; ++i) {
    objs[i].id = i;
    objs[i].xcen = ...;
    objs[i].ycen = ...;
}
```

Memory Allocation

You should be a little uneasy about **NOBJECT** in that last example. A number of questions come to mind:

- How did I know that I only had 1000 objects?
- What would I do if I had more?
- Am I wasting space if I have less?

Memory Allocation

You should be a little uneasy about **NOBJECT** in that last example. A number of questions come to mind:

- How did I know that I only had 1000 objects?
- What would I do if I had more?
- Am I wasting space if I have less?

and also

- Where is the computer putting all those IDs and positions?

Inside your job

A running program consists of a number of pieces:

Inside your job

A running program consists of a number of pieces:

- text Our instructions and read-only data

Inside your job

A running program consists of a number of pieces:

- text Our instructions and read-only data
- data Initialised global data

Inside your job

A running program consists of a number of pieces:

- text Our instructions and read-only data
- data Initialised global data
- bss Uninitialized global data

Inside your job

A running program consists of a number of pieces:

- `text` Our instructions and read-only data
- `data` Initialised global data
- `bss` Uninitialized global data
- `stack` Memory for variables in subroutines

Inside your job

A running program consists of a number of pieces:

- `text` Our instructions and read-only data
- `data` Initialised global data
- `bss` Uninitialized global data
- `stack` Memory for variables in subroutines
- `heap` Memory available to the programmer

Inside your job

A running program consists of a number of pieces:

- text Our instructions and read-only data
- data Initialised global data
- bss Uninitialized global data
- stack Memory for variables in subroutines
- heap Memory available to the programmer

All of these are mapped into a single (logical) section of RAM:

0x0	text	data	bss	heap		stack
-----	------	------	-----	------	--	-------

(0x0 is how you write a hexadecimal number in C (or C++))

Stack and Heap

It's traditional to think of the heap and stack as growing down and up respectively:

heap	↓
stack	↑

The Stack

What happens when I call a function? E.g.

```
double integrate(const float a, const float b, float (*func)(float x))  
{  
    const int nstep = 1000;           // number of steps  
    const float step = (b - a)/nstep;  
    ...  
    double ans = 0.0;  
    for (float x = a; x < b; x += step)  
        ans += func(x) * step;  
}  
  
double ans = integrate(1, 3, dfunc);
```


The Stack

What happens when I call a function? E.g.

```
double integrate(const float a, const float b, float (*func)(float)
    const int nstep = 1000;           // number of steps
    const float step = (b - a)/nstep;
    ...
    ans += func(x);
}
```

```
double ans = integrate(1, 3, dfunc);
```

We first push the values of a, b, and func onto the stack:

<dfunc>
3.0
1.0

The Stack

What happens when I call a function? E.g.

```
double integrate(const float a, const float b, float (*func)(float)
    const int nstep = 1000;           // number of steps
    const float step = (b - a)/nstep;
    ...
    ans += func(x);
}
```

```
double ans = integrate(1, 3, dfunc);
```

We first push the values of a, b, and func onto the stack:

<dfunc>
3.0
1.0

We then make space for step and nstep

???
???
<dfunc>

The Stack

We initialised `step` to 1000 and `nstep` to $(b - a)/nstep$, so:

0.002
1000
<dfunc>
3.0
1.0

The Stack

We initialised `step` to 1000 and `nstep` to $(b - a)/nstep$, so:

0.002
1000
<dfunc>
3.0
1.0

Within `integrate` we call `dfunc` with argument `x`; this pushes `x` (1.0) onto the stack:

1.0
0.002
1000
<dfunc>
3.0
1.0

The Stack

within `dfunc` we calculate `y (1.0)` and call `integrate (0, y, func)` resulting in:

0.001
1000
<func>
1.0
0.0
1.0
0.002
1000
<dfunc>
3.0
1.0

And so on.

The Stack

When `integrate` has calculated its return value, it puts it somewhere and *pops* the stack:

1.0
0.002
1000
<dfunc>
3.0
1.0

The Stack

When `integrate` has calculated its return value, it puts it somewhere and *pops* the stack:

1.0
0.002
1000
<dfunc>
3.0
1.0

`dfunc` puts *its* return value somewhere, and pops the stack:

0.002
1000
<dfunc>
3.0
1.0

The Stack

When `integrate` has calculated its return value, it puts it somewhere and *pops* the stack:

1.0
0.002
1000
<dfunc>
3.0
1.0

`dfunc` puts *its* return value somewhere, and pops the stack:

0.002
1000
<dfunc>
3.0
1.0

And finally the outer call to `integrate` finishes, saves its value, pops the stack, and we're back where we started.

Subroutine Arguments Redux

It should now be clear that:

- A language can **only** pass variables by value if it wishes to support recursion [4]
- Variables in subroutines are irretrievably lost when the routine returns
- Uninitialized values may have any value
- (Almost) all variables are stored somewhere in the process's memory

Subroutine Arguments Redux

It should now be clear that:

- A language can **only** pass variables by value if it wishes to support recursion [4]
- Variables in subroutines are irretrievably lost when the routine returns
- Uninitialized values may have any value
- (Almost) all variables are stored somewhere in the process's memory

The last bullet suggests a way to work around the first one: we can push the variable's *address* onto the stack and agree to use not the value on the stack, but the value stored at that location.

Subroutine Arguments Redux

It should now be clear that:

- A language can **only** pass variables by value if it wishes to support recursion [4]
- Variables in subroutines are irretrievably lost when the routine returns
- Uninitialized values may have any value
- (Almost) all variables are stored somewhere in the process's memory

The last bullet suggests a way to work around the first one: we can push the variable's *address* onto the stack and agree to use not the value on the stack, but the value stored at that location.

In C, `&x` is x's address, and `*px` means "use the value stored at the address px"

Pointers

A variable that holds an address is called a *pointer*; there's nothing magic about it; it just happens that you can apply the `*` operator if you want to use the value it points to.

Pointers

A variable that holds an address is called a *pointer*; there's nothing magic about it; it just happens that you can apply the `*` operator if you want to use the value it points to.

```
int i = 0;
int *pi = &i;
printf("i = %d, %d\n", i, *pi);
*pi = 10;
printf("i = %d, %d\n", i, *pi);
```

Pointers

A variable that holds an address is called a *pointer*; there's nothing magic about it; it just happens that you can apply the `*` operator if you want to use the value it points to.

```
int i = 0;
int *pi = &i;
printf("i = %d, %d\n", i, *pi);
*pi = 10;
printf("i = %d, %d\n", i, *pi);
```

If you haven't ensured that a pointer is set to a valid address, you're going to suffer. If you say

```
int *pi;
*pi = 10;
printf("i = %d\n", *pi);
```

your program may well (and should!) crash.

Arrays

By definition, given any type `type`

```
int i;  
type p[N];  
p[i] == *(p + i);  
p + i == &p[i]
```

Arrays

By definition, given any type `type`

```
int i;  
type p[N];  
p[i] == *(p + i);  
p + i == &p[i]
```

i.e. adding an integer to a pointer gives you an address larger by
`i * sizeof(type)`

Arrays

By definition, given any type `type`

```
int i;  
type p[N];  
p[i] == *(p + i);  
p + i == &p[i]
```

i.e. adding an integer to a pointer gives you an address larger by `i * sizeof(type)`

E.g. `float` is usually a 4-byte real, so if `p` is `0xffff0000`, `p + 2` is `0xffff0008`.

Arrays

By definition, given any type `type`

```
int i;  
type p[N];  
p[i] == *(p + i);  
p + i == &p[i]
```

i.e. adding an integer to a pointer gives you an address larger by `i * sizeof(type)`

E.g. `float` is usually a 4-byte real, so if `p` is `0xffff0000`, `p + 2` is `0xffff0008`.

In fact, whenever you refer to an array (`p`), it is treated as a pointer to the first element (`&p[0]`). This is why you can't pass an n-D array to a subroutine.

Strings

We introduced `char *` as a way of spelling “string” and we can now see how it works.

We can write `char str[7] = "abcdef";` [5]

Strings

We introduced `char *` as a way of spelling “string” and we can now see how it works.

We can write `char str[7] = "abcdef";` [5]

The statement `char *str = "abcdef";` is analogous to `float *x = &xx;`, and `*str` is indeed `a`.

Strings

We introduced `char *` as a way of spelling “string” and we can now see how it works.

We can write `char str[7] = "abcdef";` [5]

The statement `char *str = "abcdef";` is analogous to `float *x = &xx;`, and `*str` is indeed `a`.

The difference between `char *str = "abcdef";` and `char str[7] = "abcdef";` is where the data is actually stored; in the former case it's in the data segment, in the latter case it's on the stack.

Structs revisited

Convinced of the value of a struct such as

```
struct Object {  
    int id;           // Object IDs  
    float xcen;       // x-coordinate of center  
    float ycen;       // y-coordinate of center  
};
```

```
typedef struct Object Object;
```

I wrote a convenience function:

```
Object *newObject(const int id, const float xcen, const double ycen,  
    Object obj;  
  
    obj.id = id;  
    obj.xcen = xcen;  
    obj.ycen = ycen;  
  
    return &obj;  
}
```

Malloc

After some time spent in gdb, I remembered:

- Variables in subroutines are irretrievably lost when the routine returns

and that's exactly what this does:

```
Object *newObject(const int id, const float xcen, const double y  
    Object obj;  
    ...  
    return @\color{red}\&@obj;  
}
```

Malloc

After some time spent in gdb, I remembered:

- Variables in subroutines are irretrievably lost when the routine returns

and that's exactly what this does:

```
Object *newObject(const int id, const float xcen, const double y
    Object obj;
    ...
    return @\color{red}\&@obj;
}
```

The solution is to get a pointer to a piece of persistent memory. C provides this via a call to [malloc](#):

```
#include <stdlib.h>
```

```
Object *newObject(const int id, const float xcen, const double y
    Object @\color{red}*@obj = malloc(sizeof(Object));
    ...
    return obj;
}
```


Malloc

After some time spent in gdb, I remembered:

- Variables in subroutines are irretrievably lost when the routine returns

and that's exactly what this does:

```
Object *newObject(const int id, const float xcen, const double y
    Object obj;
    ...
    return @\color{red}\&@obj;
}
```

The solution is to get a pointer to a piece of persistent memory. C provides this via a call to [malloc](#):

```
#include <stdlib.h>
```

```
Object *newObject(const int id, const float xcen, const double y
    Object @\color{red}*@obj = malloc(sizeof(Object));
    ...
    return obj;
}
```

Pointers-to-Pointers

How about:

```
#define NOBJECT 1000
```

```
// Maximum number of obj
```

```
Object *objs[NOBJECT];
```

```
// our objects
```

Pointers-to-Pointers

How about:

```
#define NOBJECT 1000                                // Maximum number of obj
```

```
Object *objs[NOBJECT];                             // our objects
```

We know that

```
int i[10];
```

and

```
int *i;
```

are equivalent (except for the question of where the data lives).

Pointers-to-Pointers

How about:

```
#define NOBJECT 1000                                // Maximum number of obj
```

```
Object *objs[NOBJECT];                             // our objects
```

We know that

```
int i[10];
```

and

```
int *i;
```

are equivalent (except for the question of where the data lives). We can provide the needed storage with:

```
int *i = malloc(10*sizeof(int));
```

Pointers-to-Pointers

How about:

```
#define NOBJECT 1000                                // Maximum number of obj
```

```
Object *objs[NOBJECT];                             // our objects
```

We know that

```
int i[10];
```

and

```
int *i;
```

are equivalent (except for the question of where the data lives). We can provide the needed storage with:

```
int *i = malloc(10*sizeof(int));
```

So:

```
n = ...;
```

```
Object **objs = malloc(n*sizeof(Object *)); // our objects
```

Pointers-to-Pointers

How about:

```
#define NOBJECT 1000                                // Maximum number of obj
```

```
Object *objs[NOBJECT];                             // our objects
```

We know that

```
int i[10];
```

and

```
int *i;
```

are equivalent (except for the question of where the data lives). We can provide the needed storage with:

```
int *i = malloc(10*sizeof(int));
```

So:

```
n = ...;
```

```
Object **objs = malloc(n*sizeof(Object *)); // our objects
```

If you don't know n *a priori*, look up the system call [realloc](#).

Pointers-to-Pointers

How about:

```
#define NOBJECT 1000                                // Maximum number of obj
```

```
Object *objs[NOBJECT];                             // our objects
```

We know that

```
int i[10];
```

and

```
int *i;
```

are equivalent (except for the question of where the data lives). We can provide the needed storage with:

```
int *i = malloc(10*sizeof(int));
```

So:

```
n = ...;
```

```
Object **objs = malloc(n*sizeof(Object *)); // our objects
```

If you don't know n *a priori*, look up the system call [realloc](#). There is also [calloc](#) but I never use it — initialising to 0x0 is a blunt weapon.

Free

As I said, `malloc` returns persistent memory, so it's important to return it to the system when you've finished with it:

```
for (int i = 0; i != n; ++i) {  
    free(objects[i]);  
}  
free(objects);
```


Free

As I said, `malloc` returns persistent memory, so it's important to return it to the system when you've finished with it:

```
for (int i = 0; i != n; ++i) {  
    free(objects[i]);  
}  
free(objects);
```

Failure to do so results in a *memory leak*.

Free

As I said, `malloc` returns persistent memory, so it's important to return it to the system when you've finished with it:

```
for (int i = 0; i != n; ++i) {  
    free(objects[i]);  
}  
free(objects);
```

Failure to do so results in a *memory leak*.

Freeing a piece of memory more than once usually has catastrophic consequences; don't do it.

Free

As I said, `malloc` returns persistent memory, so it's important to return it to the system when you've finished with it:

```
for (int i = 0; i != n; ++i) {  
    free(objects[i]);  
}  
free(objects);
```

Failure to do so results in a *memory leak*.

Freeing a piece of memory more than once usually has catastrophic consequences; don't do it.

Another failure mode is that sometimes `malloc` can't give you the memory you want. In this case it returns 0, conventionally written `NULL`. There is not much you can do when this happens, so a reasonable response is to abort:

```
#include <stdlib.h>  
#include <assert.h>
```

```
Object *newObject(const int id, const float xcen, const double y  
    Object @\color{red}*@obj = malloc(sizeof(Object));
```

struct and typedef

I wrote

```
struct Object {  
    int id;           // Object IDs  
    ...  
};  
typedef struct Object Object;
```

struct and typedef

I wrote

```
struct Object {  
    int id;           // Object IDs  
    ...  
};
```

```
typedef struct Object Object;
```

This can be abbreviated

```
typedef struct {  
    int id;           // Object IDs  
    ...  
} Object;
```

struct and typedef

I wrote

```
struct Object {  
    int id;                // Object IDs  
    ...  
};
```

```
typedef struct Object Object;
```

This can be abbreviated

```
typedef struct {  
    int id;                // Object IDs  
    ...  
} Object;
```

There are two reasons why you might not always want to omit the name in `struct Object {...}`:

- structs can contain pointers to themselves:

```
typedef struct List {  
    struct List *prev, *next;
```

```
    ...;
```

```
} List;
```

struct and typedef

I wrote

```
struct Object {  
    int id;                // Object IDs  
    ...  
};
```

```
typedef struct Object Object;
```

This can be abbreviated

```
typedef struct {  
    int id;                // Object IDs  
    ...  
} Object;
```

There are two reasons why you might not always want to omit the name in `struct Object {...}`:

- structs can contain pointers to themselves:

```
typedef struct List {  
    struct List *prev, *next;
```

```
    ...;
```

```
} List;
```

struct and typedef

I wrote

```
struct Object {  
    int id;                // Object IDs  
    ...  
};
```

```
typedef struct Object Object;
```

This can be abbreviated

```
typedef struct {  
    int id;                // Object IDs  
    ...  
} Object;
```

There are two reasons why you might not always want to omit the name in `struct Object {...}`:

- structs can contain pointers to themselves:

```
typedef struct List {  
    struct List *prev, *next;
```

```
    ...;
```

```
} List;
```


CCD Data again


You will recall a code fragment that looked like this:

```
U16 data[4096][2048];  
...  
U16 const peak = data[y][x];           // the (x, y)th pixel  
...  
void debias(U16 data[][2048], const int nrow, const int ncol);
```

We are now in a position to do better.

CCD Data again

We need a 2-D `ncol*nrow` array of type `U16`



CCD Data again

We need a 2-D `ncol*nrow` array of type U16

Create an 1-D array `rows` of type U16 * and dimension `nrow`

CCD Data again

We need a 2-D `ncol*nrow` array of type `U16`

Create an 1-D array `rows` of type `U16 *` and dimension `nrow`

Make each element of `rows` point to the start of a row, e.g. `rows[1]` points to the first pixel in the second row of data. Then `rows[1][2]` is the value of the (2, 1) pixel.

CCD Data again

```
#include <stdlib.h>
#include <stdint.h>
#include <assert.h>

typedef uint16_t Pixel_t;

struct Image {
    Pixel_t **rows;
    int nrow, ncol;
};

typedef struct Image Image;

Image *newImage(const int ncol, int const nrow) {
    Image *im = malloc(sizeof(Image));    // the Image
    assert (im != NULL);

    im->rows = malloc(nrow*sizeof(Pixel_t *)); // pointers to r
    assert (im->rows != NULL);
```

CCD Data again

```
#include <stdlib.h>
#include <stdint.h>
#include <assert.h>

typedef uint16_t Pixel_t;

struct Image {
    Pixel_t **rows;
    int nrow, ncol;
};

typedef struct Image Image;

Image *newImage(const int ncol, int const nrow) {
    Image *im = malloc(sizeof(Image));    // the Image
    assert (im != NULL);

    im->rows = malloc(nrow*sizeof(Pixel_t *)); // pointers to r
    assert (im->rows != NULL);
```

Matrices

Now we know enough to write ourselves a matrix library

```
typedef struct Matrix {  
    float **data;  
    int nrow;  
    int ncol;  
};
```

```
Matrix *newMatrix(int nrow, int ncol);
```

```
...
```

```
Matrix *A = newMatrix(10, 10);
```

```
Matrix *B = newMatrix(10, 10);
```

Matrices

Now we know enough to write ourselves a matrix library

```
typedef struct Matrix {  
    float **data;  
    int nrow;  
    int ncol;  
};
```

```
Matrix *newMatrix(int nrow, int ncol);
```

```
...
```

```
Matrix *A = newMatrix(10, 10);
```

```
Matrix *B = newMatrix(10, 10);
```

However, you can't write:

```
Matrix *sum = A + B;
```


Matrices

Now we know enough to write ourselves a matrix library

```
typedef struct Matrix {  
    float **data;  
    int nrow;  
    int ncol;  
};
```

```
Matrix *newMatrix(int nrow, int ncol);
```

```
...
```

```
Matrix *A = newMatrix(10, 10);
```

```
Matrix *B = newMatrix(10, 10);
```

However, you can't write:

```
Matrix *sum = A + B;
```

... not until you switch to C++.

[3] In C++ you'd probably use a **const** variable but scoping rules are different in C, so a macro is appropriate

[4] this isn't quite true; it could pass a limited number of variables by using registers